

```
In [6]: # 05_second_model_edsr.ipynb
from torch.utils.data import Dataset, DataLoader

import os
from pathlib import Path

import torch
from torch import nn
from torch.utils.data import DataLoader, Subset
from torchvision import transforms
from PIL import Image
import matplotlib.pyplot as plt

print("PyTorch:", torch.__version__)

# 自动识别项目根目录(你现在在 notebooks/ 里)
PROJECT_ROOT = Path.cwd().parent if Path.cwd().name == "notebooks" else Path.cwd()
DATA_DIR = PROJECT_ROOT / "data" / "celeba"
IMG_DIR = DATA_DIR / "img_align_celeba"

print("PROJECT_ROOT:", PROJECT_ROOT)
print("IMG_DIR:", IMG_DIR, "exists:", IMG_DIR.exists())

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print("Device:", device)
```

```
PyTorch: 2.6.0+cu118
PROJECT_ROOT: D:\SuperResolution_Project-main
IMG_DIR: D:\SuperResolution_Project-main\data\celeba\img_align_celeba exists: True
Device: cuda
```

```
In [7]: # 数据集定义: HR 128x128, LR 64x64

HR_SIZE = 128
SCALE = 2
LR_SIZE = HR_SIZE // SCALE

hr_tf = transforms.Compose([
    transforms.CenterCrop(178),
    transforms.Resize((HR_SIZE, HR_SIZE), interpolation=Image.BICUBIC),
    transforms.ToTensor(),
])

to_pil = transforms.ToPILImage()

lr_from_hr = transforms.Compose([
    transforms.Resize((LR_SIZE, LR_SIZE), interpolation=Image.BICUBIC),
    transforms.ToTensor(),
])

class CelebASR(Dataset):
    def __init__(self, img_dir: Path):
        self.paths = sorted(img_dir.glob("*.jpg"))
        if len(self.paths) == 0:
            raise RuntimeError(f"No images found in {img_dir}")
    def __len__(self):
        return len(self.paths)
    def __getitem__(self, i):
```

```

        pil = Image.open(self.paths[i]).convert("RGB")
        hr = hr_tf(pil)           # [3, 128, 128]
        lr = lr_from_hr(to_pil(hr)) # [3, 64, 64]
        return lr, hr

sr_ds = CelebASR(IMG_DIR)
print("Total images:", len(sr_ds))

```

Total images: 202599

```

In [8]: from torch.utils.data import Subset, DataLoader
import torch

N = len(sr_ds)
g = torch.Generator().manual_seed(42)
perm = torch.randperm(N, generator=g)

test_ratio, val_ratio = 0.05, 0.05
test_size = int(N * test_ratio)
val_size = int(N * val_ratio)
train_size = N - val_size - test_size

train_idx = perm[:train_size]
val_idx = perm[train_size:train_size+val_size]
test_idx = perm[train_size+val_size:]

train_ds = Subset(sr_ds, train_idx.tolist())
val_ds = Subset(sr_ds, val_idx.tolist())
test_ds = Subset(sr_ds, test_idx.tolist())

print(f"Split sizes ► train: {len(train_ds)} val: {len(val_ds)} test: {len(test_ds)}")

batch_size = 16 # Mac 上可以先用 8 或 16, 慢一点但能跑
num_workers = 0 # 如果报错就改成 0

train_loader = DataLoader(train_ds, batch_size=batch_size, shuffle=True,
                         num_workers=num_workers, pin_memory=False)
val_loader = DataLoader(val_ds, batch_size=batch_size, shuffle=False,
                         num_workers=num_workers, pin_memory=False)
test_loader = DataLoader(test_ds, batch_size=batch_size, shuffle=False,
                         num_workers=num_workers, pin_memory=False)

# 简单检查一个 batch
lr_batch, hr_batch = next(iter(train_loader))
print("LR batch:", lr_batch.shape, "HR batch:", hr_batch.shape)

```

Split sizes ► train: 182341 val: 10129 test: 10129  
 LR batch: torch.Size([16, 3, 64, 64]) HR batch: torch.Size([16, 3, 128, 128])

```

In [9]: # EDSR model for 2x super-resolution

import torch
from torch import nn
import torch.nn.functional as F

class ResidualBlockNoBN(nn.Module):
    def __init__(self, n_feats=64, res_scale=0.1):
        super().__init__()
        self.res_scale = res_scale
        self.body = nn.Sequential(
            nn.Conv2d(n_feats, n_feats, kernel_size=3, padding=1),

```

```

        nn.ReLU(inplace=True),
        nn.Conv2d(n_feats, n_feats, kernel_size=3, padding=1),
    )

    def forward(self, x):
        res = self.body(x)
        return x + res * self.res_scale

class EDSR_x2(nn.Module):
    def __init__(self, in_channels=3, out_channels=3,
                 n_feats=64, n_resblocks=8, scale=2):
        super().__init__()

        # head
        self.head = nn.Conv2d(in_channels, n_feats, kernel_size=3, padding=1)

        # body: a stack of residual blocks
        blocks = [ResidualBlockNoBN(n_feats=n_feats, res_scale=0.1)
                  for _ in range(n_resblocks)]
        blocks.append(nn.Conv2d(n_feats, n_feats, kernel_size=3, padding=1))
        self.body = nn.Sequential(*blocks)

        # upsampling x2 (PixelShuffle)
        modules_upsample = [
            nn.Conv2d(n_feats, n_feats * (scale ** 2), kernel_size=3, padding=1),
            nn.PixelShuffle(scale),
            nn.Conv2d(n_feats, out_channels, kernel_size=3, padding=1),
        ]
        self.tail = nn.Sequential(*modules_upsample)

    def forward(self, x):
        x = self.head(x)
        res = self.body(x)
        x = x + res           # Long skip connection
        x = self.tail(x)
        return x

edsr = EDSR_x2().to(device)
print(edsr)

```

```
EDSR_x2(
    (head): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (body): Sequential(
        (0): ResidualBlockNoBN(
            (body): Sequential(
                (0): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
                (1): ReLU(inplace=True)
                (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
            )
        )
        (1): ResidualBlockNoBN(
            (body): Sequential(
                (0): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
                (1): ReLU(inplace=True)
                (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
            )
        )
        (2): ResidualBlockNoBN(
            (body): Sequential(
                (0): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
                (1): ReLU(inplace=True)
                (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
            )
        )
        (3): ResidualBlockNoBN(
            (body): Sequential(
                (0): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
                (1): ReLU(inplace=True)
                (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
            )
        )
        (4): ResidualBlockNoBN(
            (body): Sequential(
                (0): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
                (1): ReLU(inplace=True)
                (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
            )
        )
        (5): ResidualBlockNoBN(
            (body): Sequential(
                (0): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
                (1): ReLU(inplace=True)
                (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
            )
        )
        (6): ResidualBlockNoBN(
            (body): Sequential(
                (0): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
                (1): ReLU(inplace=True)
                (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
            )
        )
        (7): ResidualBlockNoBN(
            (body): Sequential(
                (0): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
                (1): ReLU(inplace=True)
                (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
            )
        )
        (8): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
```

```

)
(tail): Sequential(
(0): Conv2d(64, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(1): PixelShuffle(upscale_factor=2)
(2): Conv2d(64, 3, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
)
)
)

```

```
In [10]: from torch.optim import Adam

# 训练配置
num_epochs = 15           # 先用 2 epoch 测试流程; 觉得 OK 再改大一点
lr = 1e-4
log_interval = 200         # 每多少个 batch 打印一次

criterion = nn.L1Loss()    # EDSR 经典用 L1 损失
optimizer = Adam(edsr.parameters(), lr=lr)

def train_one_epoch(model, loader, optimizer, criterion, device):
    model.train()
    running_loss = 0.0

    for batch_idx, (lr_img, hr_img) in enumerate(loader):
        lr_img = lr_img.to(device)
        hr_img = hr_img.to(device)

        sr = model(lr_img)
        loss = criterion(sr, hr_img)

        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

        running_loss += loss.item()

        if (batch_idx + 1) % log_interval == 0:
            print(f" [batch {batch_idx+1:5d}] loss = {running_loss / log_interval}")
            running_loss = 0.0

@torch.no_grad()
def eval_psnr(model, loader, device, max_batches=50):
    model.eval()
    mse_sum = 0.0
    n = 0

    for b_idx, (lr_img, hr_img) in enumerate(loader):
        lr_img = lr_img.to(device)
        hr_img = hr_img.to(device)

        sr = model(lr_img).clamp(0.0, 1.0)
        mse = F.mse_loss(sr, hr_img, reduction="mean").item()
        mse_sum += mse
        n += 1

        if max_batches is not None and (b_idx + 1) >= max_batches:
            break

    avg_mse = mse_sum / max(1, n)
    psnr = -10 * torch.log10(torch.tensor(avg_mse)).item()

```

```
return psnr

# ===== 实际训练 =====
best_psnr = -1

for epoch in range(1, num_epochs + 1):
    print(f"\n==== Epoch {epoch}/{num_epochs} ====")
    train_one_epoch(edsr, train_loader, optimizer, criterion, device)
    psnr_val = eval_psnr(edsr, val_loader, device, max_batches=50)
    print(f" → Val PSNR (first 50 batches): {psnr_val:.2f} dB")

    if psnr_val > best_psnr:
        best_psnr = psnr_val
        torch.save(edsr.state_dict(), PROJECT_ROOT / "experiments" / "edsr_x2_be"
        print(" ✓ Best model updated.")
```

```
===== Epoch 1/15 =====
[batch  200] loss = 0.0576
[batch  400] loss = 0.0227
[batch  600] loss = 0.0183
[batch  800] loss = 0.0167
[batch 1000] loss = 0.0163
[batch 1200] loss = 0.0163
[batch 1400] loss = 0.0156
[batch 1600] loss = 0.0154
[batch 1800] loss = 0.0149
[batch 2000] loss = 0.0151
[batch 2200] loss = 0.0149
[batch 2400] loss = 0.0146
[batch 2600] loss = 0.0144
[batch 2800] loss = 0.0144
[batch 3000] loss = 0.0142
[batch 3200] loss = 0.0140
[batch 3400] loss = 0.0138
[batch 3600] loss = 0.0138
[batch 3800] loss = 0.0139
[batch 4000] loss = 0.0135
[batch 4200] loss = 0.0135
[batch 4400] loss = 0.0137
[batch 4600] loss = 0.0131
[batch 4800] loss = 0.0134
[batch 5000] loss = 0.0131
[batch 5200] loss = 0.0128
[batch 5400] loss = 0.0129
[batch 5600] loss = 0.0129
[batch 5800] loss = 0.0127
[batch 6000] loss = 0.0127
[batch 6200] loss = 0.0126
[batch 6400] loss = 0.0123
[batch 6600] loss = 0.0125
[batch 6800] loss = 0.0124
[batch 7000] loss = 0.0123
[batch 7200] loss = 0.0123
[batch 7400] loss = 0.0123
[batch 7600] loss = 0.0122
[batch 7800] loss = 0.0121
[batch 8000] loss = 0.0121
[batch 8200] loss = 0.0120
[batch 8400] loss = 0.0121
[batch 8600] loss = 0.0121
[batch 8800] loss = 0.0118
[batch 9000] loss = 0.0119
[batch 9200] loss = 0.0119
[batch 9400] loss = 0.0117
[batch 9600] loss = 0.0119
[batch 9800] loss = 0.0120
[batch 10000] loss = 0.0117
[batch 10200] loss = 0.0116
[batch 10400] loss = 0.0117
[batch 10600] loss = 0.0117
[batch 10800] loss = 0.0116
[batch 11000] loss = 0.0116
[batch 11200] loss = 0.0117
→ Val PSNR (first 50 batches): 33.93 dB
✓ Best model updated.
```

```
===== Epoch 2/15 =====
[batch  200] loss = 0.0114
[batch  400] loss = 0.0114
[batch  600] loss = 0.0116
[batch  800] loss = 0.0116
[batch 1000] loss = 0.0115
[batch 1200] loss = 0.0114
[batch 1400] loss = 0.0115
[batch 1600] loss = 0.0117
[batch 1800] loss = 0.0114
[batch 2000] loss = 0.0113
[batch 2200] loss = 0.0115
[batch 2400] loss = 0.0114
[batch 2600] loss = 0.0112
[batch 2800] loss = 0.0113
[batch 3000] loss = 0.0113
[batch 3200] loss = 0.0114
[batch 3400] loss = 0.0114
[batch 3600] loss = 0.0114
[batch 3800] loss = 0.0112
[batch 4000] loss = 0.0113
[batch 4200] loss = 0.0113
[batch 4400] loss = 0.0113
[batch 4600] loss = 0.0112
[batch 4800] loss = 0.0112
[batch 5000] loss = 0.0113
[batch 5200] loss = 0.0112
[batch 5400] loss = 0.0111
[batch 5600] loss = 0.0112
[batch 5800] loss = 0.0110
[batch 6000] loss = 0.0112
[batch 6200] loss = 0.0113
[batch 6400] loss = 0.0112
[batch 6600] loss = 0.0112
[batch 6800] loss = 0.0112
[batch 7000] loss = 0.0111
[batch 7200] loss = 0.0111
[batch 7400] loss = 0.0111
[batch 7600] loss = 0.0113
[batch 7800] loss = 0.0111
[batch 8000] loss = 0.0109
[batch 8200] loss = 0.0111
[batch 8400] loss = 0.0112
[batch 8600] loss = 0.0110
[batch 8800] loss = 0.0109
[batch 9000] loss = 0.0110
[batch 9200] loss = 0.0111
[batch 9400] loss = 0.0109
[batch 9600] loss = 0.0110
[batch 9800] loss = 0.0111
[batch 10000] loss = 0.0109
[batch 10200] loss = 0.0110
[batch 10400] loss = 0.0108
[batch 10600] loss = 0.0110
[batch 10800] loss = 0.0111
[batch 11000] loss = 0.0109
[batch 11200] loss = 0.0111
→ Val PSNR (first 50 batches): 34.39 dB
✓ Best model updated.
```

```
===== Epoch 3/15 =====
[batch  200] loss = 0.0109
[batch  400] loss = 0.0109
[batch  600] loss = 0.0110
[batch  800] loss = 0.0108
[batch 1000] loss = 0.0109
[batch 1200] loss = 0.0108
[batch 1400] loss = 0.0107
[batch 1600] loss = 0.0111
[batch 1800] loss = 0.0108
[batch 2000] loss = 0.0108
[batch 2200] loss = 0.0109
[batch 2400] loss = 0.0110
[batch 2600] loss = 0.0108
[batch 2800] loss = 0.0109
[batch 3000] loss = 0.0107
[batch 3200] loss = 0.0108
[batch 3400] loss = 0.0109
[batch 3600] loss = 0.0108
[batch 3800] loss = 0.0109
[batch 4000] loss = 0.0109
[batch 4200] loss = 0.0108
[batch 4400] loss = 0.0108
[batch 4600] loss = 0.0108
[batch 4800] loss = 0.0108
[batch 5000] loss = 0.0108
[batch 5200] loss = 0.0107
[batch 5400] loss = 0.0107
[batch 5600] loss = 0.0107
[batch 5800] loss = 0.0108
[batch 6000] loss = 0.0107
[batch 6200] loss = 0.0107
[batch 6400] loss = 0.0109
[batch 6600] loss = 0.0108
[batch 6800] loss = 0.0106
[batch 7000] loss = 0.0108
[batch 7200] loss = 0.0108
[batch 7400] loss = 0.0108
[batch 7600] loss = 0.0107
[batch 7800] loss = 0.0107
[batch 8000] loss = 0.0109
[batch 8200] loss = 0.0107
[batch 8400] loss = 0.0106
[batch 8600] loss = 0.0107
[batch 8800] loss = 0.0107
[batch 9000] loss = 0.0107
[batch 9200] loss = 0.0107
[batch 9400] loss = 0.0107
[batch 9600] loss = 0.0106
[batch 9800] loss = 0.0107
[batch 10000] loss = 0.0109
[batch 10200] loss = 0.0106
[batch 10400] loss = 0.0107
[batch 10600] loss = 0.0107
[batch 10800] loss = 0.0107
[batch 11000] loss = 0.0106
[batch 11200] loss = 0.0106
→ Val PSNR (first 50 batches): 34.62 dB
✓ Best model updated.
```

```
===== Epoch 4/15 =====
[batch  200] loss = 0.0106
[batch  400] loss = 0.0106
[batch  600] loss = 0.0107
[batch  800] loss = 0.0105
[batch 1000] loss = 0.0105
[batch 1200] loss = 0.0107
[batch 1400] loss = 0.0107
[batch 1600] loss = 0.0106
[batch 1800] loss = 0.0106
[batch 2000] loss = 0.0106
[batch 2200] loss = 0.0104
[batch 2400] loss = 0.0106
[batch 2600] loss = 0.0105
[batch 2800] loss = 0.0106
[batch 3000] loss = 0.0105
[batch 3200] loss = 0.0106
[batch 3400] loss = 0.0106
[batch 3600] loss = 0.0106
[batch 3800] loss = 0.0104
[batch 4000] loss = 0.0106
[batch 4200] loss = 0.0105
[batch 4400] loss = 0.0105
[batch 4600] loss = 0.0106
[batch 4800] loss = 0.0106
[batch 5000] loss = 0.0106
[batch 5200] loss = 0.0105
[batch 5400] loss = 0.0106
[batch 5600] loss = 0.0107
[batch 5800] loss = 0.0107
[batch 6000] loss = 0.0103
[batch 6200] loss = 0.0105
[batch 6400] loss = 0.0104
[batch 6600] loss = 0.0105
[batch 6800] loss = 0.0105
[batch 7000] loss = 0.0105
[batch 7200] loss = 0.0107
[batch 7400] loss = 0.0105
[batch 7600] loss = 0.0104
[batch 7800] loss = 0.0104
[batch 8000] loss = 0.0105
[batch 8200] loss = 0.0106
[batch 8400] loss = 0.0105
[batch 8600] loss = 0.0105
[batch 8800] loss = 0.0106
[batch 9000] loss = 0.0105
[batch 9200] loss = 0.0105
[batch 9400] loss = 0.0105
[batch 9600] loss = 0.0104
[batch 9800] loss = 0.0104
[batch 10000] loss = 0.0105
[batch 10200] loss = 0.0105
[batch 10400] loss = 0.0105
[batch 10600] loss = 0.0106
[batch 10800] loss = 0.0104
[batch 11000] loss = 0.0103
[batch 11200] loss = 0.0106
→ Val PSNR (first 50 batches): 34.77 dB
✓ Best model updated.
```

```
===== Epoch 5/15 =====
[batch  200] loss = 0.0104
[batch  400] loss = 0.0105
[batch  600] loss = 0.0106
[batch  800] loss = 0.0106
[batch 1000] loss = 0.0104
[batch 1200] loss = 0.0104
[batch 1400] loss = 0.0104
[batch 1600] loss = 0.0104
[batch 1800] loss = 0.0103
[batch 2000] loss = 0.0104
[batch 2200] loss = 0.0106
[batch 2400] loss = 0.0104
[batch 2600] loss = 0.0103
[batch 2800] loss = 0.0103
[batch 3000] loss = 0.0105
[batch 3200] loss = 0.0103
[batch 3400] loss = 0.0104
[batch 3600] loss = 0.0104
[batch 3800] loss = 0.0104
[batch 4000] loss = 0.0104
[batch 4200] loss = 0.0104
[batch 4400] loss = 0.0105
[batch 4600] loss = 0.0104
[batch 4800] loss = 0.0103
[batch 5000] loss = 0.0105
[batch 5200] loss = 0.0104
[batch 5400] loss = 0.0104
[batch 5600] loss = 0.0104
[batch 5800] loss = 0.0103
[batch 6000] loss = 0.0103
[batch 6200] loss = 0.0104
[batch 6400] loss = 0.0104
[batch 6600] loss = 0.0103
[batch 6800] loss = 0.0103
[batch 7000] loss = 0.0103
[batch 7200] loss = 0.0103
[batch 7400] loss = 0.0103
[batch 7600] loss = 0.0104
[batch 7800] loss = 0.0104
[batch 8000] loss = 0.0104
[batch 8200] loss = 0.0103
[batch 8400] loss = 0.0105
[batch 8600] loss = 0.0105
[batch 8800] loss = 0.0103
[batch 9000] loss = 0.0104
[batch 9200] loss = 0.0104
[batch 9400] loss = 0.0103
[batch 9600] loss = 0.0104
[batch 9800] loss = 0.0102
[batch 10000] loss = 0.0102
[batch 10200] loss = 0.0102
[batch 10400] loss = 0.0104
[batch 10600] loss = 0.0103
[batch 10800] loss = 0.0104
[batch 11000] loss = 0.0104
[batch 11200] loss = 0.0102
→ Val PSNR (first 50 batches): 34.87 dB
✓ Best model updated.
```

```
===== Epoch 6/15 =====
[batch  200] loss = 0.0103
[batch  400] loss = 0.0103
[batch  600] loss = 0.0102
[batch  800] loss = 0.0104
[batch 1000] loss = 0.0104
[batch 1200] loss = 0.0104
[batch 1400] loss = 0.0102
[batch 1600] loss = 0.0103
[batch 1800] loss = 0.0103
[batch 2000] loss = 0.0103
[batch 2200] loss = 0.0103
[batch 2400] loss = 0.0104
[batch 2600] loss = 0.0103
[batch 2800] loss = 0.0102
[batch 3000] loss = 0.0103
[batch 3200] loss = 0.0104
[batch 3400] loss = 0.0103
[batch 3600] loss = 0.0101
[batch 3800] loss = 0.0104
[batch 4000] loss = 0.0102
[batch 4200] loss = 0.0102
[batch 4400] loss = 0.0103
[batch 4600] loss = 0.0102
[batch 4800] loss = 0.0103
[batch 5000] loss = 0.0104
[batch 5200] loss = 0.0103
[batch 5400] loss = 0.0102
[batch 5600] loss = 0.0102
[batch 5800] loss = 0.0102
[batch 6000] loss = 0.0103
[batch 6200] loss = 0.0102
[batch 6400] loss = 0.0101
[batch 6600] loss = 0.0102
[batch 6800] loss = 0.0102
[batch 7000] loss = 0.0104
[batch 7200] loss = 0.0102
[batch 7400] loss = 0.0103
[batch 7600] loss = 0.0102
[batch 7800] loss = 0.0103
[batch 8000] loss = 0.0103
[batch 8200] loss = 0.0102
[batch 8400] loss = 0.0102
[batch 8600] loss = 0.0103
[batch 8800] loss = 0.0102
[batch 9000] loss = 0.0103
[batch 9200] loss = 0.0101
[batch 9400] loss = 0.0101
[batch 9600] loss = 0.0102
[batch 9800] loss = 0.0103
[batch 10000] loss = 0.0101
[batch 10200] loss = 0.0103
[batch 10400] loss = 0.0102
[batch 10600] loss = 0.0102
[batch 10800] loss = 0.0101
[batch 11000] loss = 0.0102
[batch 11200] loss = 0.0103
→ Val PSNR (first 50 batches): 34.93 dB
✓ Best model updated.
```

```
===== Epoch 7/15 =====
[batch  200] loss = 0.0102
[batch  400] loss = 0.0101
[batch  600] loss = 0.0102
[batch  800] loss = 0.0101
[batch 1000] loss = 0.0102
[batch 1200] loss = 0.0102
[batch 1400] loss = 0.0102
[batch 1600] loss = 0.0102
[batch 1800] loss = 0.0103
[batch 2000] loss = 0.0101
[batch 2200] loss = 0.0102
[batch 2400] loss = 0.0101
[batch 2600] loss = 0.0101
[batch 2800] loss = 0.0102
[batch 3000] loss = 0.0102
[batch 3200] loss = 0.0101
[batch 3400] loss = 0.0102
[batch 3600] loss = 0.0102
[batch 3800] loss = 0.0103
[batch 4000] loss = 0.0101
[batch 4200] loss = 0.0104
[batch 4400] loss = 0.0101
[batch 4600] loss = 0.0102
[batch 4800] loss = 0.0102
[batch 5000] loss = 0.0101
[batch 5200] loss = 0.0101
[batch 5400] loss = 0.0101
[batch 5600] loss = 0.0101
[batch 5800] loss = 0.0102
[batch 6000] loss = 0.0101
[batch 6200] loss = 0.0102
[batch 6400] loss = 0.0101
[batch 6600] loss = 0.0102
[batch 6800] loss = 0.0102
[batch 7000] loss = 0.0101
[batch 7200] loss = 0.0103
[batch 7400] loss = 0.0102
[batch 7600] loss = 0.0101
[batch 7800] loss = 0.0103
[batch 8000] loss = 0.0102
[batch 8200] loss = 0.0100
[batch 8400] loss = 0.0101
[batch 8600] loss = 0.0102
[batch 8800] loss = 0.0103
[batch 9000] loss = 0.0102
[batch 9200] loss = 0.0101
[batch 9400] loss = 0.0101
[batch 9600] loss = 0.0101
[batch 9800] loss = 0.0102
[batch 10000] loss = 0.0102
[batch 10200] loss = 0.0102
[batch 10400] loss = 0.0102
[batch 10600] loss = 0.0101
[batch 10800] loss = 0.0101
[batch 11000] loss = 0.0100
[batch 11200] loss = 0.0101
→ Val PSNR (first 50 batches): 35.01 dB
✓ Best model updated.
```

```
===== Epoch 8/15 =====
[batch  200] loss = 0.0102
[batch  400] loss = 0.0101
[batch  600] loss = 0.0101
[batch  800] loss = 0.0102
[batch 1000] loss = 0.0101
[batch 1200] loss = 0.0102
[batch 1400] loss = 0.0101
[batch 1600] loss = 0.0101
[batch 1800] loss = 0.0101
[batch 2000] loss = 0.0100
[batch 2200] loss = 0.0102
[batch 2400] loss = 0.0102
[batch 2600] loss = 0.0101
[batch 2800] loss = 0.0102
[batch 3000] loss = 0.0101
[batch 3200] loss = 0.0101
[batch 3400] loss = 0.0103
[batch 3600] loss = 0.0101
[batch 3800] loss = 0.0100
[batch 4000] loss = 0.0101
[batch 4200] loss = 0.0102
[batch 4400] loss = 0.0100
[batch 4600] loss = 0.0101
[batch 4800] loss = 0.0101
[batch 5000] loss = 0.0100
[batch 5200] loss = 0.0100
[batch 5400] loss = 0.0099
[batch 5600] loss = 0.0101
[batch 5800] loss = 0.0101
[batch 6000] loss = 0.0101
[batch 6200] loss = 0.0102
[batch 6400] loss = 0.0101
[batch 6600] loss = 0.0100
[batch 6800] loss = 0.0100
[batch 7000] loss = 0.0102
[batch 7200] loss = 0.0102
[batch 7400] loss = 0.0102
[batch 7600] loss = 0.0101
[batch 7800] loss = 0.0101
[batch 8000] loss = 0.0102
[batch 8200] loss = 0.0102
[batch 8400] loss = 0.0101
[batch 8600] loss = 0.0100
[batch 8800] loss = 0.0101
[batch 9000] loss = 0.0101
[batch 9200] loss = 0.0101
[batch 9400] loss = 0.0101
[batch 9600] loss = 0.0101
[batch 9800] loss = 0.0101
[batch 10000] loss = 0.0101
[batch 10200] loss = 0.0100
[batch 10400] loss = 0.0102
[batch 10600] loss = 0.0101
[batch 10800] loss = 0.0100
[batch 11000] loss = 0.0101
[batch 11200] loss = 0.0100
→ Val PSNR (first 50 batches): 35.06 dB
✓ Best model updated.
```

```
===== Epoch 9/15 =====
[batch  200] loss = 0.0101
[batch  400] loss = 0.0101
[batch  600] loss = 0.0101
[batch  800] loss = 0.0101
[batch 1000] loss = 0.0100
[batch 1200] loss = 0.0100
[batch 1400] loss = 0.0100
[batch 1600] loss = 0.0100
[batch 1800] loss = 0.0100
[batch 2000] loss = 0.0101
[batch 2200] loss = 0.0100
[batch 2400] loss = 0.0101
[batch 2600] loss = 0.0102
[batch 2800] loss = 0.0100
[batch 3000] loss = 0.0101
[batch 3200] loss = 0.0101
[batch 3400] loss = 0.0099
[batch 3600] loss = 0.0101
[batch 3800] loss = 0.0100
[batch 4000] loss = 0.0100
[batch 4200] loss = 0.0100
[batch 4400] loss = 0.0100
[batch 4600] loss = 0.0101
[batch 4800] loss = 0.0101
[batch 5000] loss = 0.0100
[batch 5200] loss = 0.0101
[batch 5400] loss = 0.0101
[batch 5600] loss = 0.0101
[batch 5800] loss = 0.0100
[batch 6000] loss = 0.0100
[batch 6200] loss = 0.0101
[batch 6400] loss = 0.0101
[batch 6600] loss = 0.0100
[batch 6800] loss = 0.0101
[batch 7000] loss = 0.0100
[batch 7200] loss = 0.0100
[batch 7400] loss = 0.0099
[batch 7600] loss = 0.0101
[batch 7800] loss = 0.0101
[batch 8000] loss = 0.0100
[batch 8200] loss = 0.0101
[batch 8400] loss = 0.0100
[batch 8600] loss = 0.0100
[batch 8800] loss = 0.0102
[batch 9000] loss = 0.0100
[batch 9200] loss = 0.0099
[batch 9400] loss = 0.0100
[batch 9600] loss = 0.0099
[batch 9800] loss = 0.0099
[batch 10000] loss = 0.0101
[batch 10200] loss = 0.0100
[batch 10400] loss = 0.0100
[batch 10600] loss = 0.0099
[batch 10800] loss = 0.0102
[batch 11000] loss = 0.0100
[batch 11200] loss = 0.0101
→ Val PSNR (first 50 batches): 35.08 dB
✓ Best model updated.
```

```
===== Epoch 10/15 =====
[batch  200] loss = 0.0100
[batch  400] loss = 0.0100
[batch  600] loss = 0.0100
[batch  800] loss = 0.0099
[batch 1000] loss = 0.0101
[batch 1200] loss = 0.0100
[batch 1400] loss = 0.0101
[batch 1600] loss = 0.0101
[batch 1800] loss = 0.0101
[batch 2000] loss = 0.0099
[batch 2200] loss = 0.0099
[batch 2400] loss = 0.0100
[batch 2600] loss = 0.0101
[batch 2800] loss = 0.0100
[batch 3000] loss = 0.0099
[batch 3200] loss = 0.0100
[batch 3400] loss = 0.0100
[batch 3600] loss = 0.0101
[batch 3800] loss = 0.0099
[batch 4000] loss = 0.0101
[batch 4200] loss = 0.0100
[batch 4400] loss = 0.0100
[batch 4600] loss = 0.0102
[batch 4800] loss = 0.0101
[batch 5000] loss = 0.0100
[batch 5200] loss = 0.0099
[batch 5400] loss = 0.0100
[batch 5600] loss = 0.0100
[batch 5800] loss = 0.0100
[batch 6000] loss = 0.0100
[batch 6200] loss = 0.0100
[batch 6400] loss = 0.0100
[batch 6600] loss = 0.0099
[batch 6800] loss = 0.0100
[batch 7000] loss = 0.0101
[batch 7200] loss = 0.0099
[batch 7400] loss = 0.0100
[batch 7600] loss = 0.0100
[batch 7800] loss = 0.0100
[batch 8000] loss = 0.0099
[batch 8200] loss = 0.0099
[batch 8400] loss = 0.0099
[batch 8600] loss = 0.0099
[batch 8800] loss = 0.0100
[batch 9000] loss = 0.0100
[batch 9200] loss = 0.0099
[batch 9400] loss = 0.0100
[batch 9600] loss = 0.0100
[batch 9800] loss = 0.0099
[batch 10000] loss = 0.0100
[batch 10200] loss = 0.0100
[batch 10400] loss = 0.0100
```

```

-----
KeyboardInterrupt                               Traceback (most recent call last)

Cell In[10], line 61
  59 for epoch in range(1, num_epochs + 1):
  60     print(f"\n==== Epoch {epoch}/{num_epochs} ====")
---> 61     train_one_epoch(edsr, train_loader, optimizer, criterion, device)
  62     psnr_val = eval_psnr(edsr, val_loader, device, max_batches=50)
  63     print(f" → Val PSNR (first 50 batches): {psnr_val:.2f} dB")

Cell In[10], line 26, in train_one_epoch(model, loader, optimizer, criterion, device)
  23 loss.backward()
  24 optimizer.step()
---> 26 running_loss += loss.item()
  28 if (batch_idx + 1) % log_interval == 0:
  29     print(f" [batch {batch_idx+1:5d}] loss = {running_loss / log_interval:.4f}")

KeyboardInterrupt:

```

```

In [11]: import torch
         from pathlib import Path

         ckpt_path = Path("outputs/checkpoints/edsr_quicksave.pth")
         ckpt_path.parent.mkdir(parents=True, exist_ok=True)

         torch.save(edsr.state_dict(), ckpt_path)
         print("✓ EDSR checkpoint saved to:", ckpt_path)

✓ EDSR checkpoint saved to: outputs\checkpoints\edsr_quicksave.pth

```

```

In [14]: import torch
         import torch.nn.functional as F
         import matplotlib.pyplot as plt

         @torch.no_grad()
         def show_sr_examples(model, loader, device, num_images=4):
             model.eval()

             # 取一个 batch
             lr_batch, hr_batch = next(iter(loader))
             lr_batch = lr_batch.to(device)
             hr_batch = hr_batch.to(device)

             # 前向得到 SR, 并先搬回 CPU
             sr_batch = model(lr_batch).clamp(0.0, 1.0).cpu()

             # HR 也搬回 CPU, 后面画图用
             hr_batch = hr_batch.cpu()

             fig, axes = plt.subplots(num_images, 3, figsize=(9, 3 * num_images))
             if num_images == 1:
                 axes = [axes]

             for i in range(num_images):
                 # LR 这里最终也要在 CPU 上画图
                 lr_up = F.interpolate(
                     lr_batch[i].unsqueeze(0),
                     size=hr_batch.shape[-2:],
                     mode="nearest"

```

```
    ).cpu().squeeze(0)

    hr_img = hr_batch[i]      # 已经在 CPU
    sr_img = sr_batch[i]      # 已经在 CPU

    axes[i][0].imshow(lr_up.permute(1, 2, 0))
    axes[i][0].set_title("LR (upsampled)")
    axes[i][0].axis("off")

    axes[i][1].imshow(sr_img.permute(1, 2, 0))
    axes[i][1].set_title("EDSR SR")
    axes[i][1].axis("off")

    axes[i][2].imshow(hr_img.permute(1, 2, 0))
    axes[i][2].set_title("HR")
    axes[i][2].axis("off")

    plt.tight_layout()
    plt.show()

# 调用
show_sr_examples(edsr, test_loader, device, num_images=4)
```



In [ ]: